

Progressive Autonomy as Preference Learning

A Formalization of Trust Calibration for Agentic Tool Use

Changkun Ou

March 4, 2026

Abstract

We formalize trust calibration for agentic tool use (deciding when an automated agent’s proposed action may execute autonomously versus require human approval) as a preference-learning problem. A policy gateway maintains a Gaussian-process posterior over a latent human risk-tolerance function, observed through a probit likelihood on binary approve/deny feedback, and escalates to the human exactly where the approval outcome is most uncertain. We show this is structurally an instance of Preferential Bayesian Optimization, inheriting its inference machinery (approximate Gaussian-process classification) and its sample-efficiency argument (uncertainty-targeted querying), while differing in objective: classifying an action space into allow/block/ask regions rather than optimizing a design.

1 Background and Related Work

Deciding how much autonomy to delegate to an automated system is a classical human-in-the-loop control problem; the technical core here, recovering a human’s latent acceptability function from sparse binary feedback, is *preference learning*. Chu and Ghahramani [5] introduced Gaussian-process preference learning, placing a Gaussian-process (GP) prior over a latent utility and linking observed human choices to it through a probit likelihood. We adopt exactly this construction, specialized to unary approve/deny feedback. The same latent-utility-with-probit model underlies preference-driven sequential decision making: González et al. [7] formalized *Preferential Bayesian Optimization* (PBO), embedding GP preference learning in the query loop of Bayesian optimization [2, 15]. Our policy gateway is structurally an instance of this framework, differing only in objective (classifying an action space rather than optimizing a design), as Table 1 makes precise.

The inferential machinery is classical GP classification: a GP prior, a non-Gaussian probit likelihood, and an analytically intractable posterior approximated by the Laplace method or Expectation Propagation [9], developed comprehensively by Rasmussen and Williams [12]. Treating the human query budget as a scarce resource makes the ASK region an *acquisition* rule in the sense of active learning [14] and Bayesian optimization [15]: interruptions are spent where the expected information about the allow/block boundary is greatest. Finally, drifting risk tolerance is a non-stationarity problem; we model it with a time-decaying kernel component in the spirit of non-stationary covariance functions [11] and, for abrupt shifts, Bayesian online changepoint detection [1].

This progressive view of autonomy has deep roots in the trust-in-automation literature: Lee and See [8] characterized appropriate reliance as the alignment of trust with actual system trust-worthiness, and de Visser et al. [6] extended this to *longitudinal* trust calibration in human–robot teams, precisely the dynamic our time-decaying kernel (§6) is meant to capture. The concern has resurfaced sharply for large-language-model agents, where graduated autonomy has emerged as an

explicit deployment axis alongside raw capability [10]. Recent work catalogs the risks of increasingly agentic systems [3], argues for visibility and oversight mechanisms over deployed agents [4], and proposes governance practices in which a human retains approval authority over consequential actions [16]. These accounts are largely qualitative and taxonomic: they argue *why* graduated autonomy matters and *what* should be governed, but leave the escalation policy itself as a fixed, hand-specified tier. Our contribution is the missing mechanism: a learning rule that lets the auto-approve/escalate boundary adapt from human feedback rather than being set by hand.

2 Setup

At each decision point $t = 1, 2, \dots$, the policy gateway observes a proposed agent action $a_t \in \mathcal{A}$ and an execution context $c_t \in \mathcal{C}$, where:

$$a_t = (\text{tool_name}, \text{args}, \text{target_resource}), \quad (1)$$

$$c_t = (\text{repo_state}, \text{task_desc}, \text{session_history}). \quad (2)$$

A human supervisor provides binary feedback $y_t \in \{0, 1\}$ (deny/approve). We write $x_t := (a_t, c_t) \in \mathcal{X} = \mathcal{A} \times \mathcal{C}$ for the joint input.

3 Latent Risk Tolerance

Definition 1 (Risk tolerance function). *There exists a latent function $f: \mathcal{X} \rightarrow \mathbb{R}$ encoding the human’s risk tolerance, such that approval probability follows a probit observation model:*

$$\Pr(y = 1 \mid x) = \Phi(f(x)), \quad (3)$$

where Φ is the standard normal CDF.

Remark 1. *This is structurally identical to the observation model in Preferential Bayesian Optimization [7] and, more fundamentally, in Gaussian-process preference learning [5]. In standard PBO, the human expresses a preference between two candidates (x_i, x_j) via $\Pr(x_i \succ x_j) = \Phi(f(x_i) - f(x_j))$. Here, the comparison is against an implicit internal threshold: the human approves when $f(x)$ exceeds their acceptable-risk boundary. The unary case is a degenerate pairwise comparison against a fixed reference point.*

4 Gaussian Process Prior and Posterior

Place a GP prior over f [12]:

$$f \sim \mathcal{GP}(\mu_0, k(x, x')). \quad (4)$$

The kernel k decomposes over the input structure. A natural choice is a product kernel:

$$k(x, x') = k_{\text{tool}}(a, a') \cdot k_{\text{ctx}}(c, c') \cdot k_{\text{time}}(t, t'), \quad (5)$$

where:

- k_{tool} encodes similarity between actions (e.g., shared tool name, overlapping argument patterns, same reversibility class),
- k_{ctx} captures context similarity (same repository, file type, task category),
- k_{time} handles non-stationarity (see §6).

After observing $\mathcal{D}_N = \{(x_t, y_t)\}_{t=1}^N$, the posterior is:

$$p(f \mid \mathcal{D}_N) \propto \mathcal{GP}(\mu_0, k) \cdot \prod_{t=1}^N \Phi(f(x_t))^{y_t} (1 - \Phi(f(x_t)))^{1-y_t}. \quad (6)$$

This is analytically intractable due to the non-Gaussian likelihood. Approximate inference proceeds via the Laplace approximation [12] or Expectation Propagation [9], the same machinery used in PBO.

5 The Policy Gateway Decision Rule

Given the posterior predictive distribution at a new point x_* :

$$\hat{p}(x_*) := \mathbb{E}_{f \mid \mathcal{D}_N}[\Phi(f(x_*))], \quad (7)$$

the gateway applies a three-tier decision:

$$\text{decision}(x_*) = \begin{cases} \text{ALLOW} & \text{if } \hat{p}(x_*) > \tau_{\text{high}}, \\ \text{BLOCK} & \text{if } \hat{p}(x_*) < \tau_{\text{low}}, \\ \text{ASK} & \text{otherwise.} \end{cases} \quad (8)$$

The ASK region $[\tau_{\text{low}}, \tau_{\text{high}}]$ plays the role of an *acquisition function* [14, 15]: the system queries the human precisely where the model is most uncertain about the approval outcome, maximizing the expected value of information per human interruption.

Remark 2. *The target operating point (85–90% auto-approve, 10–15% human escalation) emerges naturally as the posterior concentrates. Early in a project, most actions fall in the ASK region. As the model learns, the ASK band narrows, and the auto-approve rate rises without any manual threshold tuning.*

6 Non-Stationarity

Human risk tolerance drifts: early in a project the supervisor is cautious; as trust accumulates for familiar patterns, they become permissive. Model this via a time-decayed kernel component:

$$k_{\text{time}}(t, t') = \exp\left(-\frac{|t - t'|}{\lambda}\right), \quad (9)$$

where $\lambda > 0$ is a lengthscale controlling the forgetting rate. This gives recent approve/deny signals more weight, a principled analog of the intuition that “agents earn trust over time.”

For computational efficiency, an equivalent effect can be achieved through a sliding window of the most recent W observations, or via online changepoint detection when the supervisor’s behavior shifts abruptly (e.g., moving to a new codebase).

7 Correlated Generalization

A key advantage over a naïve contextual bandit (which treats each (a, c) independently) is *correlated generalization* through the kernel. Concretely:

- Approving `write_file` to `/workspace/src/` transfers evidence to `write_file` to `/workspace/test/`, since k_{tool} and k_{ctx} assign high similarity.
- Denying `execute_sql` with a `DROP` argument propagates caution to `execute_sql` with `TRUNCATE`, without the human having to deny each variant individually.
- A new tool with no interaction history inherits the prior μ_0 , which maps to `ASK`, the fail-safe default.

8 Connection to PBO

The mapping between the trust calibration problem and Preferential Bayesian Optimization [7] is summarized in Table 1.

Component	PBO (Optimization)	Trust Calibration (Policy)
Input space $\mathcal{X}$	Design parameters	(action, context) pairs
Latent function f	Objective to maximize	Risk tolerance to learn
Human feedback	Pairwise preference $x_i \succ x_j$	Unary approve/deny
Observation model	$\Phi(f(x_i) - f(x_j))$	$\Phi(f(x))$
Acquisition	Next query to evaluate	Next action to escalate
Goal	Find $x^* = \arg \max f$	Learn the ALLOW/BLOCK boundary

Table 1: Structural correspondence between PBO and trust calibration.

The “preferential” aspect is literal: the human is expressing preferences over what the agent should be allowed to do. The same mathematical machinery (GP priors, probit likelihoods, approximate posterior inference) transfers directly. The difference is the objective: PBO seeks to *optimize*, while trust calibration seeks to *classify* the action space into allow/block/ask regions with minimal human queries.

9 Datasets and Evaluation

Empirically grounding the gateway requires data pairing proposed agent actions with human approve/deny judgments. The closest public resource is R-Judge [17], which supplies multi-turn agent interaction records annotated by humans with binary safe/unsafe labels across a range of risk scenarios; it is a natural source for a cold-start prior μ_0 and for calibrating k_{tool} and k_{ctx} . Broader agent-safety benchmarks such as Agent-SafetyBench [18] and ToolEmu [13] widen the action and context coverage, but their risk labels are produced by automated judges rather than per-action human approval, so they are better suited to stress-testing the learned boundary than to fitting f itself.

A structural gap remains: no public dataset captures the *longitudinal, per-supervisor* signal that the non-stationarity model (§6) assumes. Existing corpora provide single-shot, aggregated annotations and do not track an individual supervisor’s risk tolerance drifting over the course of a project. Validating the time-decaying kernel therefore requires either a controlled user study or a simulation with deliberately drifting synthetic annotators, with R-Judge serving as the static initializer. We treat this as an inherent limitation of currently available data rather than a deficiency of the model.

References

- [1] Ryan Prescott Adams and David J. C. MacKay. Bayesian online changepoint detection. *arXiv preprint arXiv:0710.3742*, 2007. URL <https://arxiv.org/abs/0710.3742>.
- [2] Eric Brochu, Vlad M. Cora, and Nando de Freitas. A tutorial on Bayesian optimization of expensive cost functions, with application to active user modeling and hierarchical reinforcement learning. *arXiv preprint arXiv:1012.2599*, 2010. URL <https://arxiv.org/abs/1012.2599>.
- [3] Alan Chan, Rebecca Salganik, Alva Markelius, Chris Pang, Nitarshan Rajkumar, Dmitrii Krasheninnikov, Lauro Langosco, Zhonghao He, Yawen Duan, Micah Carroll, Michelle Lin, Alex Mayhew, Katherine Collins, Maryam Molamohammadi, John Burden, Wanru Zhao, Shalaleh Rismeni, Konstantinos Voudouris, Umang Bhatt, Adrian Weller, David Krueger, and Tegan Maharaj. Harms from increasingly agentic algorithmic systems. In *Proceedings of the 2023 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*, pages 651–666. ACM, 2023. doi: 10.1145/3593013.3594033. URL <https://doi.org/10.1145/3593013.3594033>.
- [4] Alan Chan, Carson Ezell, Max Kaufmann, Kevin Wei, Lewis Hammond, Herbie Bradley, Emma Bluemke, Nitarshan Rajkumar, David Krueger, Noam Kolt, Lennart Heim, and Markus Anderljung. Visibility into AI agents. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*, pages 958–973. ACM, 2024. doi: 10.1145/3630106.3658948. URL <https://doi.org/10.1145/3630106.3658948>.
- [5] Wei Chu and Zoubin Ghahramani. Preference learning with Gaussian processes. In *Proceedings of the 22nd International Conference on Machine Learning (ICML)*, pages 137–144. ACM Press, 2005. doi: 10.1145/1102351.1102369. URL <https://doi.org/10.1145/1102351.1102369>.
- [6] Ewart J. de Visser, Marieke M. M. Peeters, Malte F. Jung, Spencer Kohn, Tyler H. Shaw, Richard Pak, and Mark A. Neerinx. Towards a theory of longitudinal trust calibration in human–robot teams. *International Journal of Social Robotics*, 12(2):459–478, 2020. doi: 10.1007/s12369-019-00596-x. URL <https://doi.org/10.1007/s12369-019-00596-x>.
- [7] Javier González, Zhenwen Dai, Andreas Damianou, and Neil D. Lawrence. Preferential Bayesian optimization. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, volume 70 of *Proceedings of Machine Learning Research*, pages 1282–1291. PMLR, 2017. URL <https://proceedings.mlr.press/v70/gonzalez17a.html>.
- [8] John D. Lee and Katrina A. See. Trust in automation: Designing for appropriate reliance. *Human Factors*, 46(1):50–80, 2004. URL https://doi.org/10.1518/hfes.46.1.50_30392.
- [9] Thomas P. Minka. Expectation propagation for approximate Bayesian inference. In *Proceedings of the Seventeenth Conference on Uncertainty in Artificial Intelligence (UAI)*, pages 362–369. Morgan Kaufmann, 2001. URL <https://tminka.github.io/papers/ep/minka-ep-uai.pdf>.
- [10] Meredith Ringel Morris, Jascha Sohl-Dickstein, Noah Fiedel, Tris Warkentin, Allan Dafoe, Aleksandra Faust, Clement Farabet, and Shane Legg. Position: Levels of AGI for operationalizing progress on the path to AGI. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *Proceedings of Machine Learning Research*, pages 36308–36321. PMLR, 2024. URL <https://proceedings.mlr.press/v235/morris24b.html>.

- [11] Christopher J. Paciorek and Mark J. Schervish. Nonstationary covariance functions for Gaussian process regression. In *Advances in Neural Information Processing Systems 16 (NIPS)*. MIT Press, 2003. URL <https://proceedings.neurips.cc/paper/2003/hash/326a8c055c0d04f5b06544665d8bb3ea-Abstract.html>.
- [12] Carl Edward Rasmussen and Christopher K. I. Williams. *Gaussian Processes for Machine Learning*. MIT Press, Cambridge, MA, 2006. ISBN 978-0-262-18253-9. doi: 10.7551/mitpress/3206.001.0001. URL <https://gaussianprocess.org/gpml/chapters/RW.pdf>.
- [13] Yangjun Ruan, Honghua Dong, Andrew Wang, Silviu Pitis, Yongchao Zhou, Jimmy Ba, Yann Dubois, Chris J. Maddison, and Tatsunori Hashimoto. Identifying the risks of LM agents with an LM-emulated sandbox. In *The Twelfth International Conference on Learning Representations (ICLR)*, 2024. URL <https://openreview.net/forum?id=GEcwtMk1uA>. Spotlight.
- [14] Burr Settles. Active learning literature survey. Computer Sciences Technical Report 1648, University of Wisconsin–Madison, 2009. URL <https://research.cs.wisc.edu/techreports/2009/TR1648.pdf>.
- [15] Bobak Shahriari, Kevin Swersky, Ziyu Wang, Ryan P. Adams, and Nando de Freitas. Taking the human out of the loop: A review of Bayesian optimization. *Proceedings of the IEEE*, 104(1):148–175, 2016. doi: 10.1109/JPROC.2015.2494218. URL <https://doi.org/10.1109/JPROC.2015.2494218>.
- [16] Yonadav Shavit, Sandhini Agarwal, Miles Brundage, Steven Adler, Cullen O’Keefe, Rosie Campbell, Teddy Lee, Pamela Mishkin, Tyna Eloundou, Alan Hickey, Katarina Slama, Lama Ahmad, Paul McMillan, Alex Beutel, Alexandre Passos, and David G. Robinson. Practices for governing agentic AI systems. White paper, OpenAI, 2023. URL <https://cdn.openai.com/papers/practices-for-governing-agentic-ai-systems.pdf>.
- [17] Tongxin Yuan, Zhiwei He, Lingzhong Dong, Yiming Wang, Ruijie Zhao, Tian Xia, Lizhen Xu, Binglin Zhou, Fangqi Li, Zhuosheng Zhang, Rui Wang, and Gongshen Liu. R-judge: Benchmarking safety risk awareness for LLM agents. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 1467–1490. Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.findings-emnlp.79. URL <https://aclanthology.org/2024.findings-emnlp.79/>.
- [18] Zhexin Zhang, Shiyao Cui, Yida Lu, Jingzhuo Zhou, Junxiao Yang, Hongning Wang, and Minlie Huang. Agent-safetybench: Evaluating the safety of LLM agents, 2024. URL <https://arxiv.org/abs/2412.14470>.